

Aim

To develop a Retrieval-Augmented Generation (RAG) system using the GEMMA GROQ model for document-based question answering. The system integrates various AI tools and libraries to process, store, and retrieve relevant document information efficiently.

Objective

1. Implement a document ingestion and embedding system using Google Generative AI embeddings.
2. Create a retrieval mechanism to find relevant document chunks based on user queries.
3. Develop a user-friendly interface using Streamlit to interact with the system.
4. Integrate the GEMMA GROQ model to generate accurate answers to user queries based on the retrieved document context.

Summary

This project focuses on building a RAG system using GEMMA GROQ and various Langchain components. The system ingests documents from a specified directory, processes them into vector embeddings, and stores them using FAISS. Users can input questions through a Streamlit interface, and the system retrieves relevant document chunks to generate accurate responses using the GEMMA GROQ model.

Tools and Libraries Used

- **Streamlit:** For building the user interface.
- **Langchain:** For managing document processing and retrieval.
- **FAISS:** For storing and retrieving vector embeddings.
- **PyPDF2:** For loading PDF documents.
- **Google Generative AI:** For generating embeddings.
- **dotenv:** For loading API keys from the environment.
- **GEMMA GROQ:** For generating responses to user queries.

Procedure

1. **Setup and Configuration:**
 - Load environment variables and API keys.
 - Initialize the Streamlit interface.
2. **Document Processing:**
 - Load documents from a specified directory using PyPDFDirectoryLoader.
 - Split documents into chunks using RecursiveCharacterTextSplitter.

3. Embedding and Vector Storage:

- Generate embeddings for document chunks using GoogleGenerativeAIEmbeddings.
- Store embeddings in a FAISS vector store.

4. User Interaction:

- Provide a text input field for users to ask questions.
- Embed documents into the vector store when prompted.

5. Query Processing:

- Create a document chain using the GEMMA GROQ model and a prompt template.
- Retrieve relevant document chunks using the FAISS vector store.
- Generate and display responses based on retrieved documents.

CODE:

```
import streamlit as st

import os

from langchain_groq import ChatGroq

from langchain.text_splitter import RecursiveCharacterTextSplitter

from langchain.chains.combine_documents import create_stuff_documents_chain

from langchain_core.prompts import ChatPromptTemplate

from langchain.chains import create_retrieval_chain

from langchain_community.vectorstores import FAISS

from langchain_community.document_loaders import PyPDFDirectoryLoader

from langchain_google_genai import GoogleGenerativeAIEmbeddings

from dotenv import load_dotenv

import os

load_dotenv()

## load the GROQ And OpenAI API KEY

groq_api_key=os.getenv('GROQ_API_KEY')

os.environ["GOOGLE_API_KEY"]=os.getenv("GOOGLE_API_KEY")
```

```
st.title("Gemma Model Document Q&A")

llm=ChatGroq(groq_api_key=groq_api_key,
             model_name="Gemma-7b-it")

prompt=ChatPromptTemplate.from_template(
    """"
Answer the questions based on the provided context only.

Please provide the most accurate response based on the question

<context>
{context}
<context>

Questions:{input}

"""
)

def vector_embedding():

    if "vectors" not in st.session_state:

        st.session_state.embeddings=GoogleGenerativeAIEmbeddings(model = "models/embedding-001")
        st.session_state.loader=PyPDFDirectoryLoader("./data") ## Data Ingestion
        st.session_state.docs=st.session_state.loader.load() ## Document Loading
        st.session_state.text_splitter=RecursiveCharacterTextSplitter(chunk_size=1000,chunk_overlap=200)
## Chunk Creation
        st.session_state.final_documents=st.session_state.text_splitter.split_documents(st.session_state.docs[:20]) #splitting
        st.session_state.vectors=FAISS.from_documents(st.session_state.final_documents,st.session_state.embeddings) #vector OpenAI embeddings
```

```
prompt1=st.text_input("Enter Your Question From Documents")

if st.button("Documents Embedding"):
    vector_embedding()
    st.write("Vector Store DB Is Ready")

import time

if prompt1:
    document_chain=create_stuff_documents_chain(llm,prompt)
    retriever=st.session_state.vectors.as_retriever()
    retrieval_chain=create_retrieval_chain(retriever,document_chain)
    start=time.process_time()
    response=retrieval_chain.invoke({'input':prompt1})
    print("Response time :",time.process_time()-start)
    st.write(response['answer'])

# With a streamlit expander
with st.expander("Document Similarity Search"):
    # Find the relevant chunks
    for i, doc in enumerate(response["context"]):
        st.write(doc.page_content)
        st.write("-----")
```

Highlights

- **Integration of GEMMA GROQ Model:** The system utilizes the GEMMA GROQ model to generate precise answers based on the provided document context.
- **Vector Embeddings with Google Generative AI:** Embeddings are generated using state-of-the-art models, ensuring high-quality vector representations.
- **Efficient Document Retrieval:** FAISS is used to efficiently store and retrieve vector embeddings, facilitating quick access to relevant document chunks.
- **User-Friendly Interface:** Streamlit provides an intuitive interface for users to interact with the system, ask questions, and view responses.

Conclusion

The RAG system effectively combines advanced AI models and tools to provide accurate document-based question answering. By leveraging the power of GEMMA GROQ, Google Generative AI, and Langchain, the project demonstrates a robust solution for information retrieval and generation. The use of Streamlit ensures a seamless user experience, making the system accessible and easy to use.